

Monitoring my application like a boss



fagossa



fabgutier



fabian.gutierrez

Agenda

- Why are we here?
- Monitoring
- Logging
- Push model andThen demo
- Pull model andThen demo
- Questions

Why are we here?

**because we need to know
when the house is on fire**







Monitoring

Follow-up the state of an application

Can be achieved by any/all of these three things?

- Logging
- Tracing (this one get philosophical)
- Metrics

Logging vs Tracing

What is the difference between Tracing and Logging?



6



1

From the terminology point of view and in general, what is the difference between a 'tracing' and a 'logging' ?

Thanks!

java

.net

windows

logging

tracing

share edit flag

asked Dec 2 '14 at 8:21



pencilCake

16.8k ● 53 ● 164 ● 286

[Event Logging vs. Tracing](#) sums it up nicely. – [Filburt](#) Dec 2 '14 at 8:28

[add a comment](#)

[start a bounty](#)

Logging - Solved problem

- No more `ssh + tail -f output.log`
- **Strings** containing diverse information
 - (level, user, host, etc)
- **Slf4j** (logback, others)
- Individual records matter
- **Direct business value** (€)
- Non-ephemeral
- Logs can be **used as metrics**



Logging - examples

1

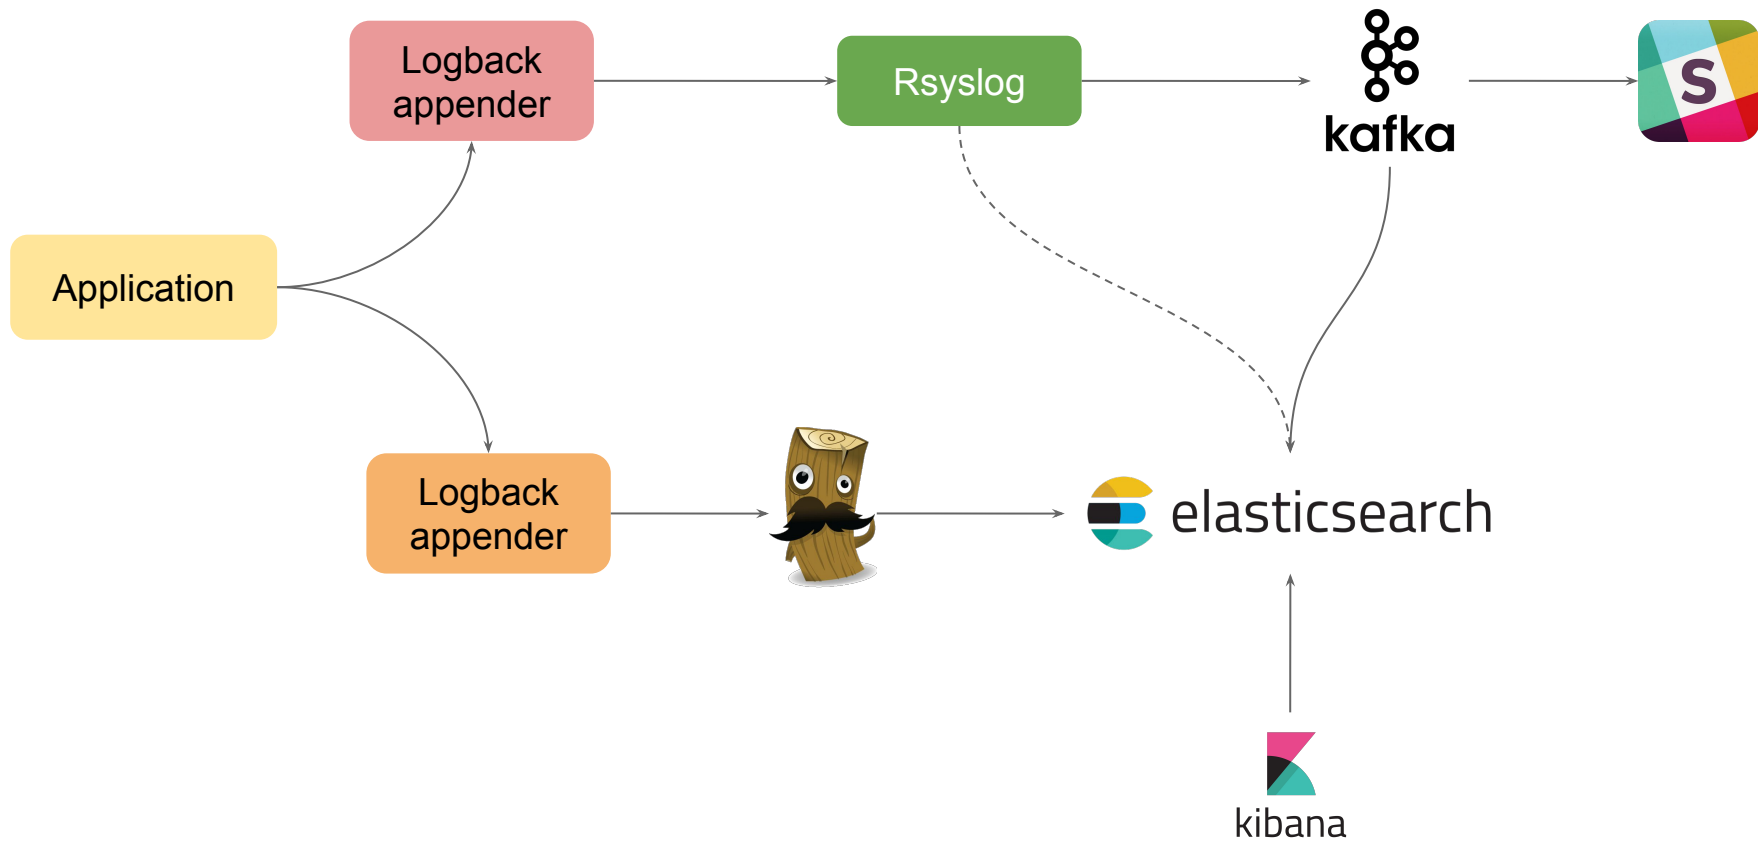
217.0.0.1 - jean [INFO] GET /products 200

2

host=217.0.0.1 **user**=jean **method**=GET **url**=/products **time**=200

3

```
{  
  "host": "217.0.0.1",  
  "user": "jean",  
  "method": "GET",  
  "url": "/products",  
  "time": 200  
}
```



Metrics

- (hopefully) Multidimensional data
- **Direct tech value**
 - Response times
 - Complex flows
- Business value (?)
- **Ephemeral**
- Easier alerting
- Logs => metrics
 - <https://github.com/google/mtail/wiki>

- **Examples**

http_request_duration_seconds_count
{

method="GET",

path="/metrics",

status="2xx"

} 4

3


current_users

3.0


4

How can I get
metrics out of
my app?






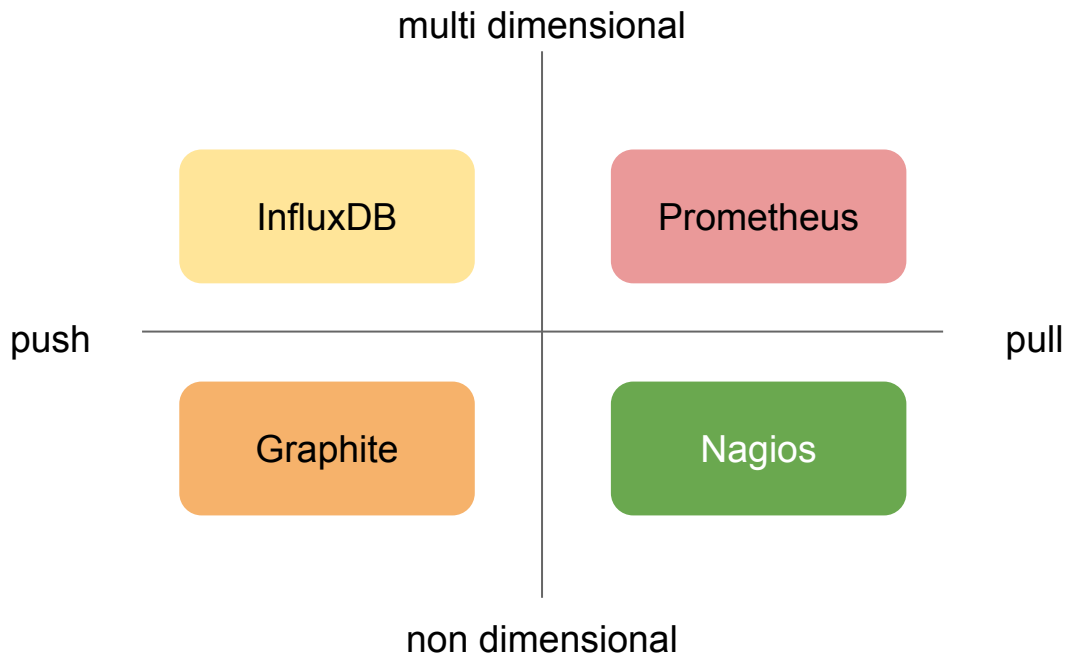
PUSH



PULL



We are talking about this

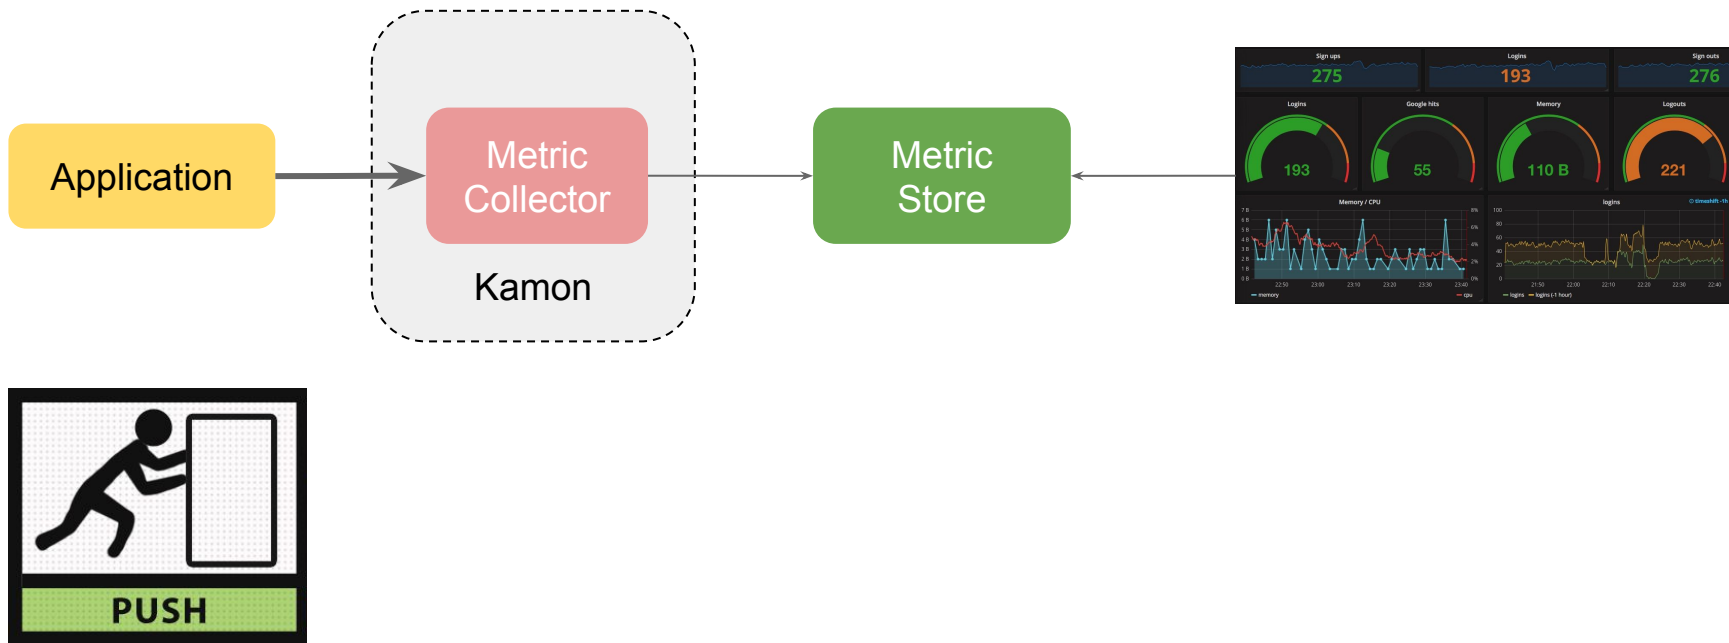


Push approach

Push those
metrics my
friend



Push approach



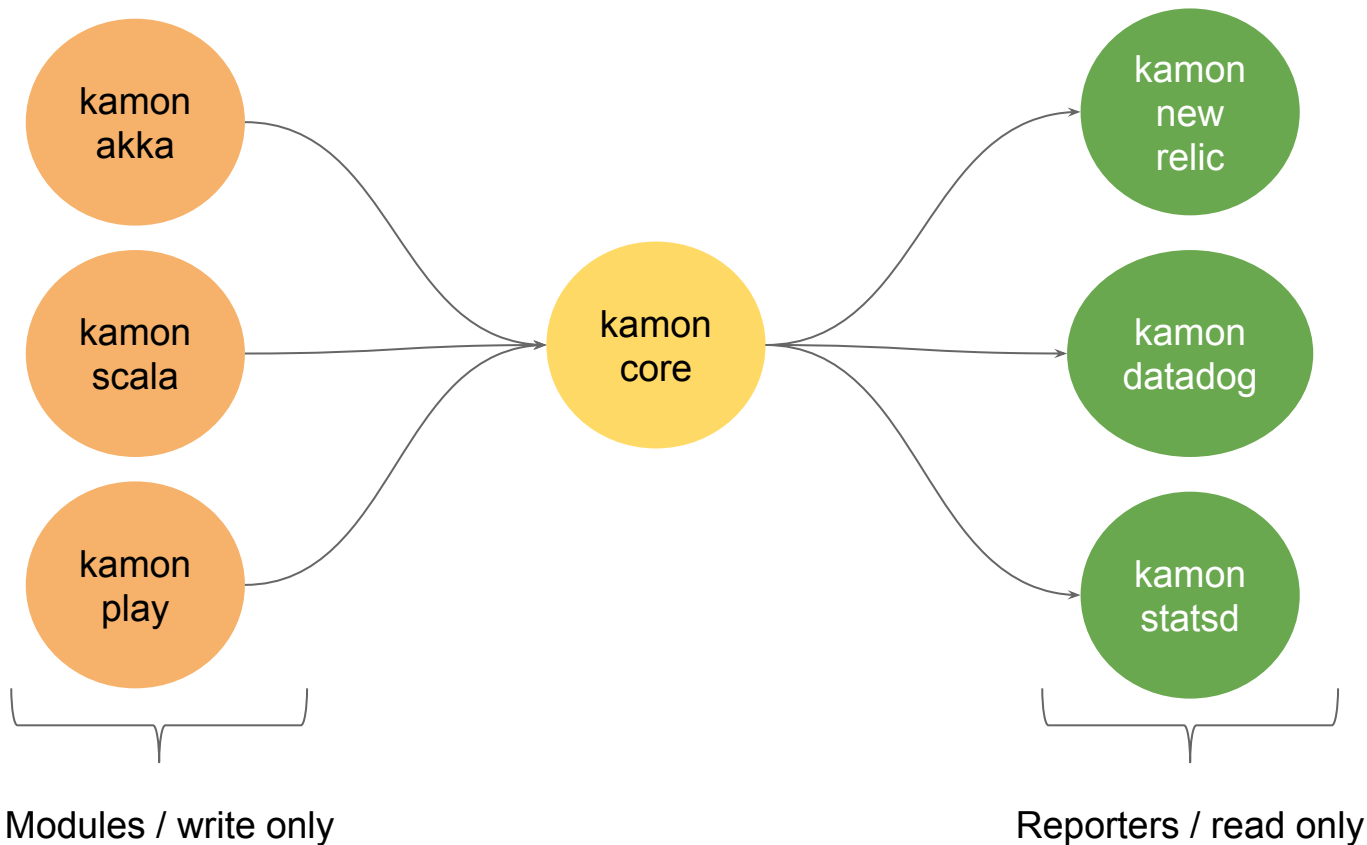
Kamon

Monitoring tool for the JVM

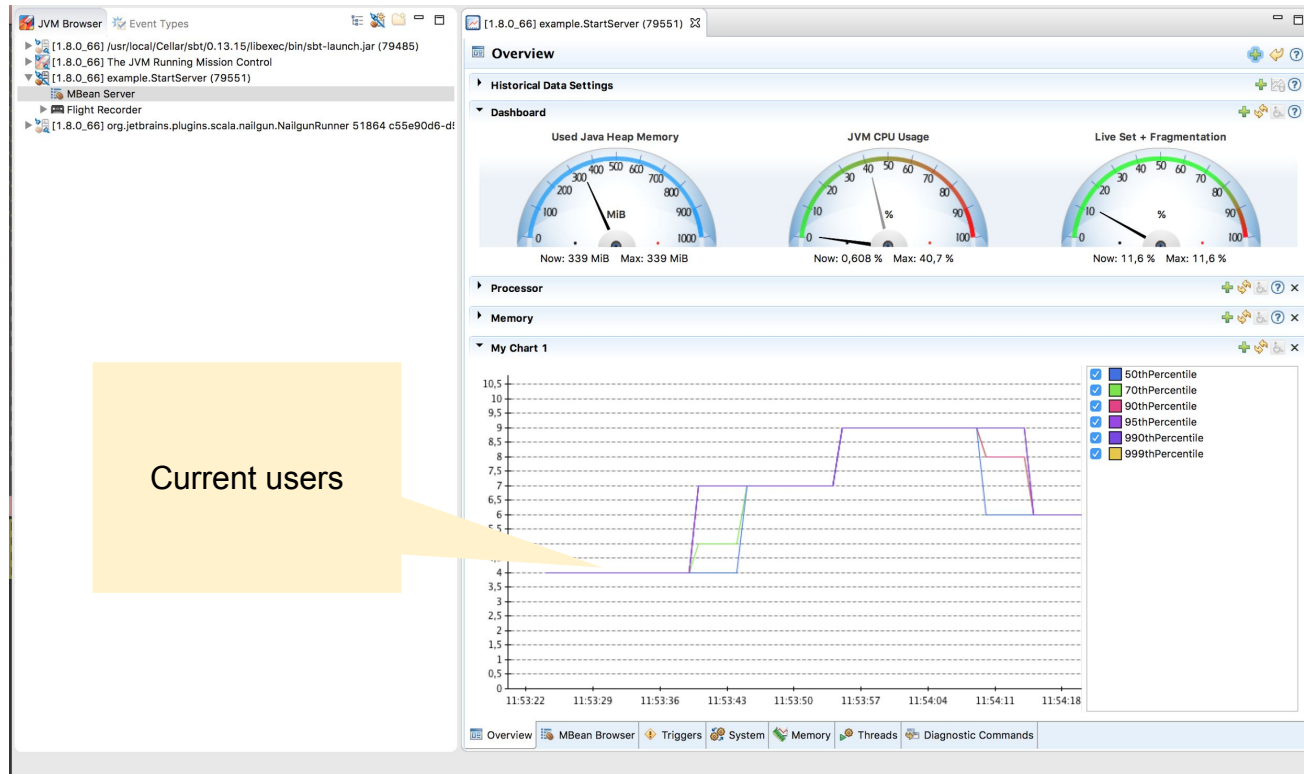

- Open source
- Metrics and tracing API
- **Instrumentation** for common libraries (akka, play, etc)
- **Collection** and **reporting** are **separate**
 - Instrument once, report anywhere



High Level view



Kamon + JMX Reporter



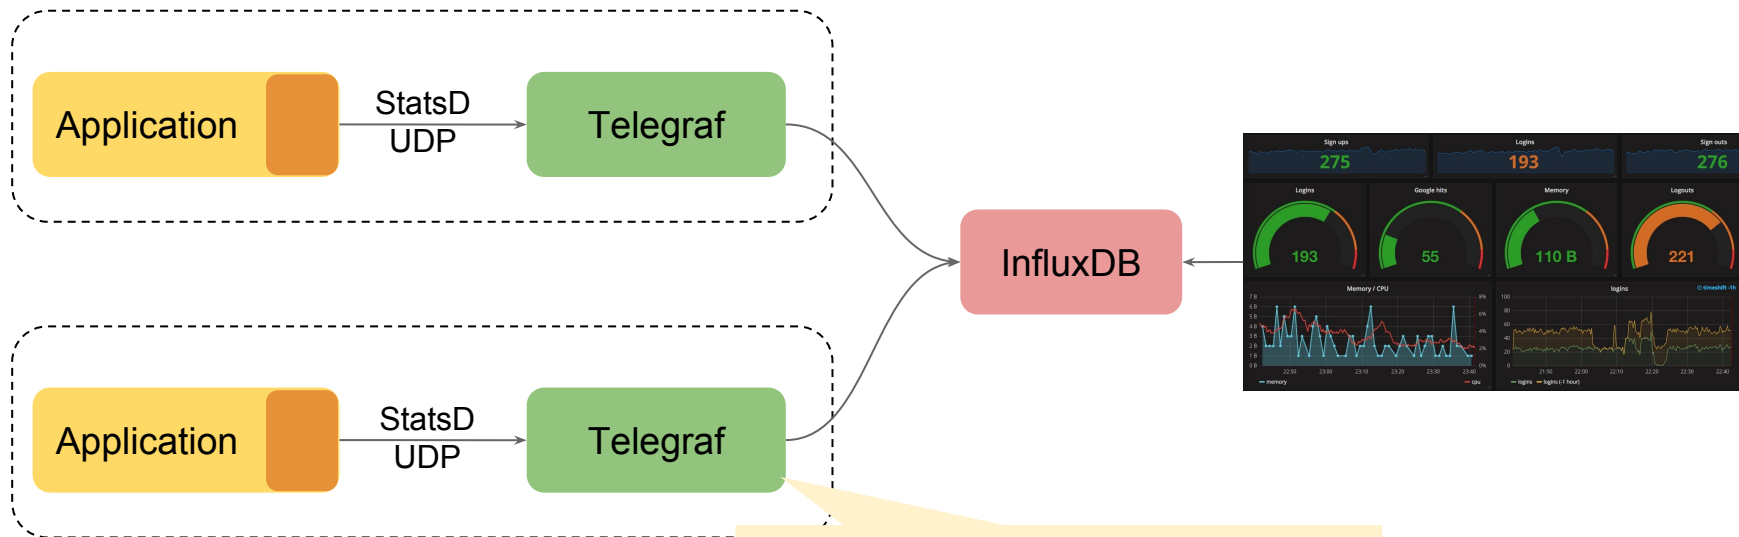
DEMO GODS



PLEASE LET THIS DEMO WORK

meme-generator.net

Kamon + Telegraf + influxDB



- Agent that accepts StatsD protocol metrics
- Aggregates and parses metrics
- Periodically forwards the metrics to InfluxDB

Kamon and Actors


For each actor you have access to **4 metrics**:

- Errors
- Mailbox-size
- Processing-time
- Time-in-mailbox



Kamon - The good

RECAP



- **Push** approach
- **Great integration** with the JVM
- Several modules (JMX, StatsD, etc)
- **Active project:** A new version (1.0.0) came out a couple of months ago



Kamon - The bad?

RECAP

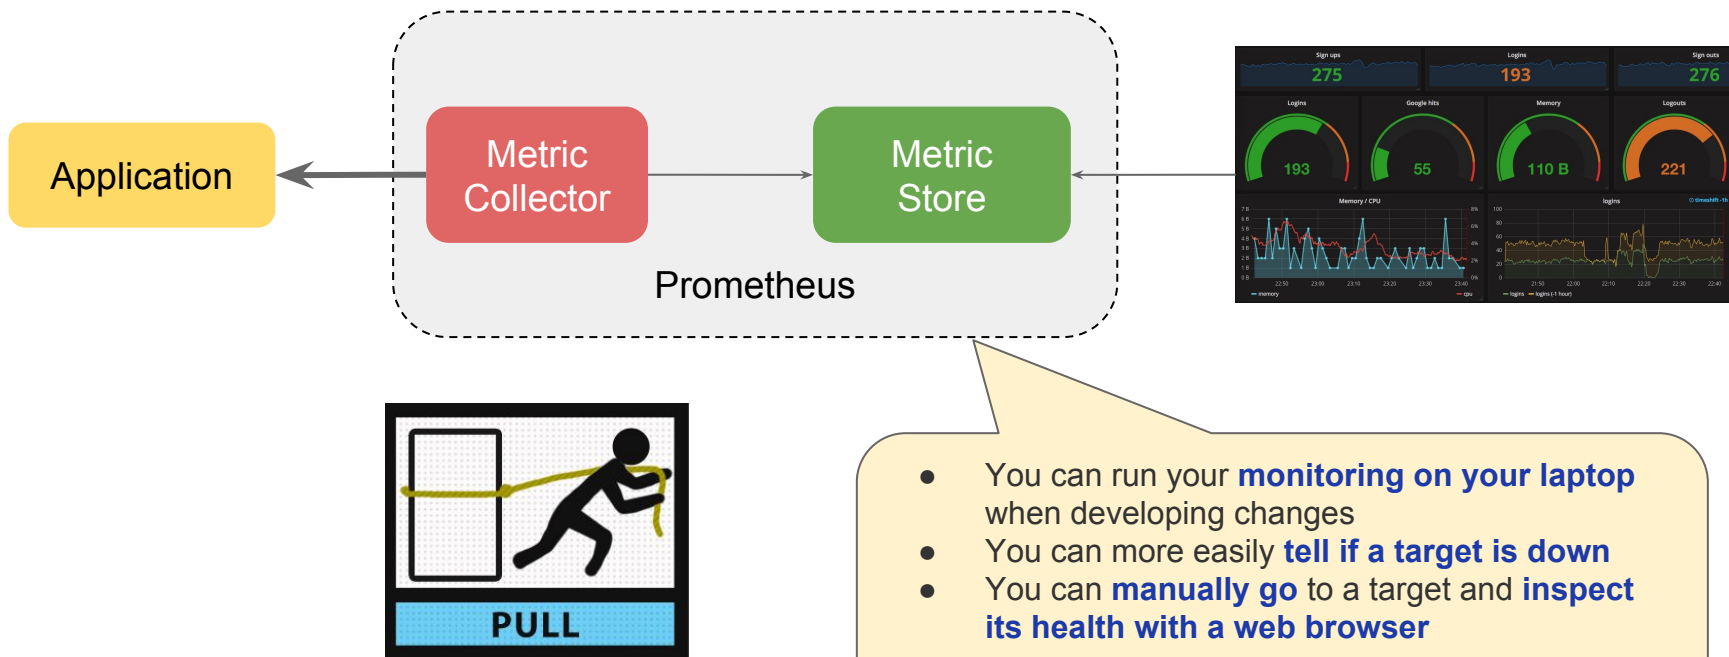


- **Bytecode** instrumentation (?)
- Working with modules is sometimes **confusing** (à la Spring)
- Potential **bytecode incompatibilities**
- Outside the JVM

Pull approach



Pull approach



Prometheus

System monitoring tool with built-in timeseries DB

- **Integrates collecting** and **reporting**
- Metric API
- **Alerting already provided**
- Only **numeric timeseries** metrics

It is not

- Don't do logging or tracing
- Not distributed storage (**only local**) by design!



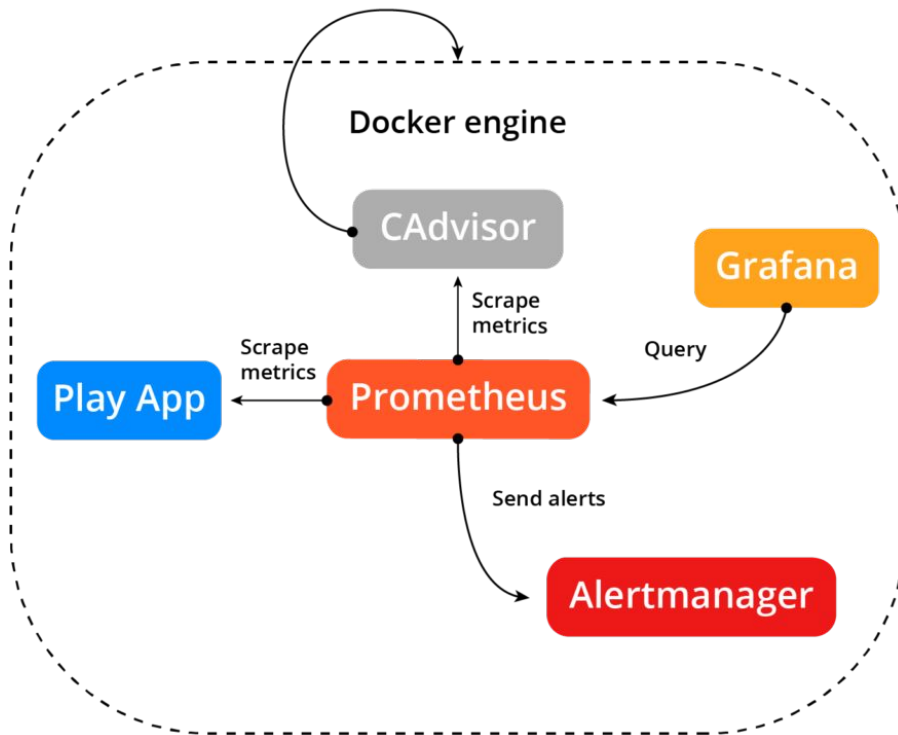
/metrics

- | | | |
|---|---|-------------|
| 1 | <code>http_request_duration_seconds_count {
 method="GET",
 path="/metrics",
 status="2xx"} -----></code> | 2 |
| 2 | <code>http_request_duration_seconds_sum {
 method="GET",
 path="/metrics",
 status="2xx"} -----></code> | 0.065599873 |
| 3 | <code>http_request_mismatch_total -----></code> | 1.0 |
| 4 | <code>play_current_users -----></code> | 3.0 |
| 5 | <code>play_requests_total -----></code> | 1.0 |

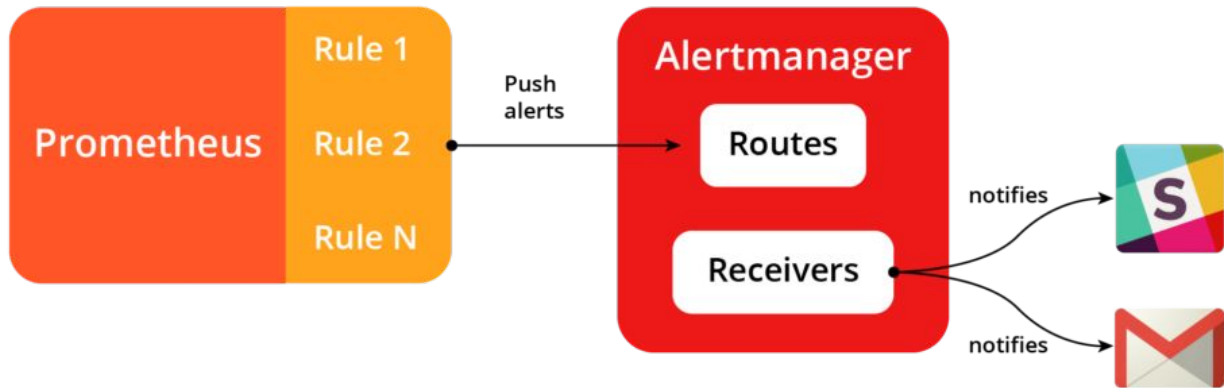
Prometheus + docker



High level view



Alerting



Alerting

1
ALERT `low_connected_users`

IF `play_current_users` < 2 2

FOR 30s 3

LABELS {

`severity` = "warning"

}

ANNOTATIONS {

4 `summary` = "Instance {{ \$labels.instance }} under lower load",

`description` =

 "{{ \$labels.instance }} of job {{ \$labels.job }} is under lower load.",

}

RULES


① _____

② _____

③ _____

④ _____

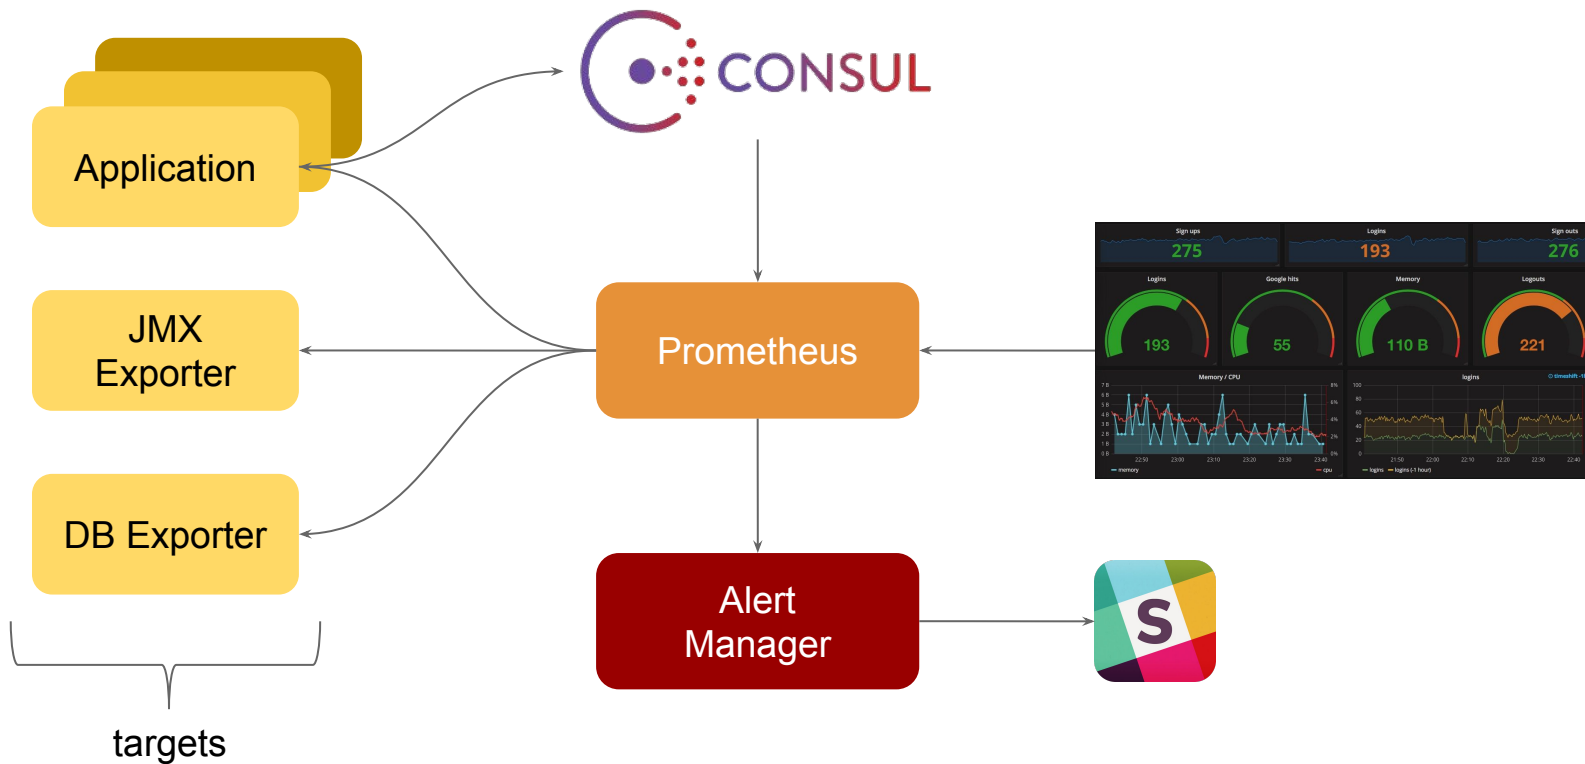
DEMO GODS



PLEASE LET THIS DEMO WORK

meme-generator.net

A more complex architecture



Prometheus - The good

RECAP



- **Pull** approach
- **Prepackaged** solution (collect + storage)
- Easy to start with
- Simple metric API
- **Active project** (version 2.0 just came out)
- Lots of **exporters**
- prometheus-akka seems nice



Prometheus - The bad?

RECAP



- **Ephemeral** persistence (?)
- What to do after a **few weeks of logs**? (existing adaptors to influxDB)
- App **overhead**?
- Kamon-prometheus bridge :(

From zero to hero

Start with high level metrics, like user experienced response time

Go a bit deeper and analyze sections of functionality within your app

Go even deeper and analyze the core components of your app

How long does a **login** take?

How long did the "select all products" **JDBC call** take?

How many **messages** is handling this **actor**?




Mixing things up

<https://github.com/lonelyplanet/prometheus-akka-http>

<https://github.com/Workday/prometheus-akka>

<https://github.com/MonsantoCo/kamon-prometheus>



Conclusions



- Both approaches are **robust** enough
- Good **integrations** for both
- **Don't guess ... monitor**
- Does not matter which approach ... **choose one**



Going further

- <https://github.com/fagossa/play-prometheus>
- <http://blog.xebia.fr/2017/07/28/superviser-mon-application-play-avec-prometheus>
- <https://en.fabernovel.com/insights/tech-en/alerting-in-prometheus-or-how-i-can-sleep-well-at-night>